



Cassandra

Data Modelling and Queries with CQL3

By Markus Klems
(2013)



HOW TO WRITE A CV




Leverage the NoSQL boom

Data Model

- Keyspace (Database)
- Column Family (Table)
- Keys and Columns

A column



the timestamp field is
used by Cassandra for conflict
resolution: “Last Write Wins”

Column family (Table)

partition key

columns ...

partition key

101	email	name	tel	
	ab@c.to	otto	12345	
103	email	name	tel	tel2
	karl@a.b	karl	6789	12233
104	name			
	linda			

Table with standard PRIMARY KEY

```
CREATE TABLE messages (  
  msg_id timeuuid PRIMARY KEY,  
  author text,  
  body text  
);
```

Table: Tweets

PRIMARY KEY
= msg_id

9990	author	body
	otto	Hello World!
9991	author	body
	linda	Hi, Otto

Table with compound PRIMARY KEY

```
CREATE TABLE timeline (  
    user_id uuid,  
    msg_id timeuuid,  
    author text,  
    body text,  
    PRIMARY KEY (user_id, msg_id)  
);
```


“Wide-row” Table: Timeline


PRIMARY KEY = user_id + msg_id

partition key column


103	msg_id	author	body
	9990	otto	Hello World!
103	msg_id	author	body
	9991	linda	Hi @otto

Timeline Table is partitioned by user and locally clustered by msg

103	msg_id	...
	9990	...
103	msg_id	...
	9994	...
104	msg_id	...
	9881	...
104	msg_id	...
	9999	...



211	msg_id	...
	8090	...
211	msg_id	...
	8555	...
211	msg_id	...
	9678	...
212	msg_id	...
	9877	...



Comparison: RDBMS vs. Cassandra

RDBMS Data Design

Users Table

user_id	name	email
101	otto	o@t.to

Tweets Table

tweet_id	author_id	body
9990	101	Hello!

Followers Table

id	follows_id	followed_id
4321	104	101

Cassandra Data Design

Users Table

user_id	name	email
101	otto	o@t.to

Tweets Table

tweet_id	author_id	name	body
9990	101	otto	Hello!

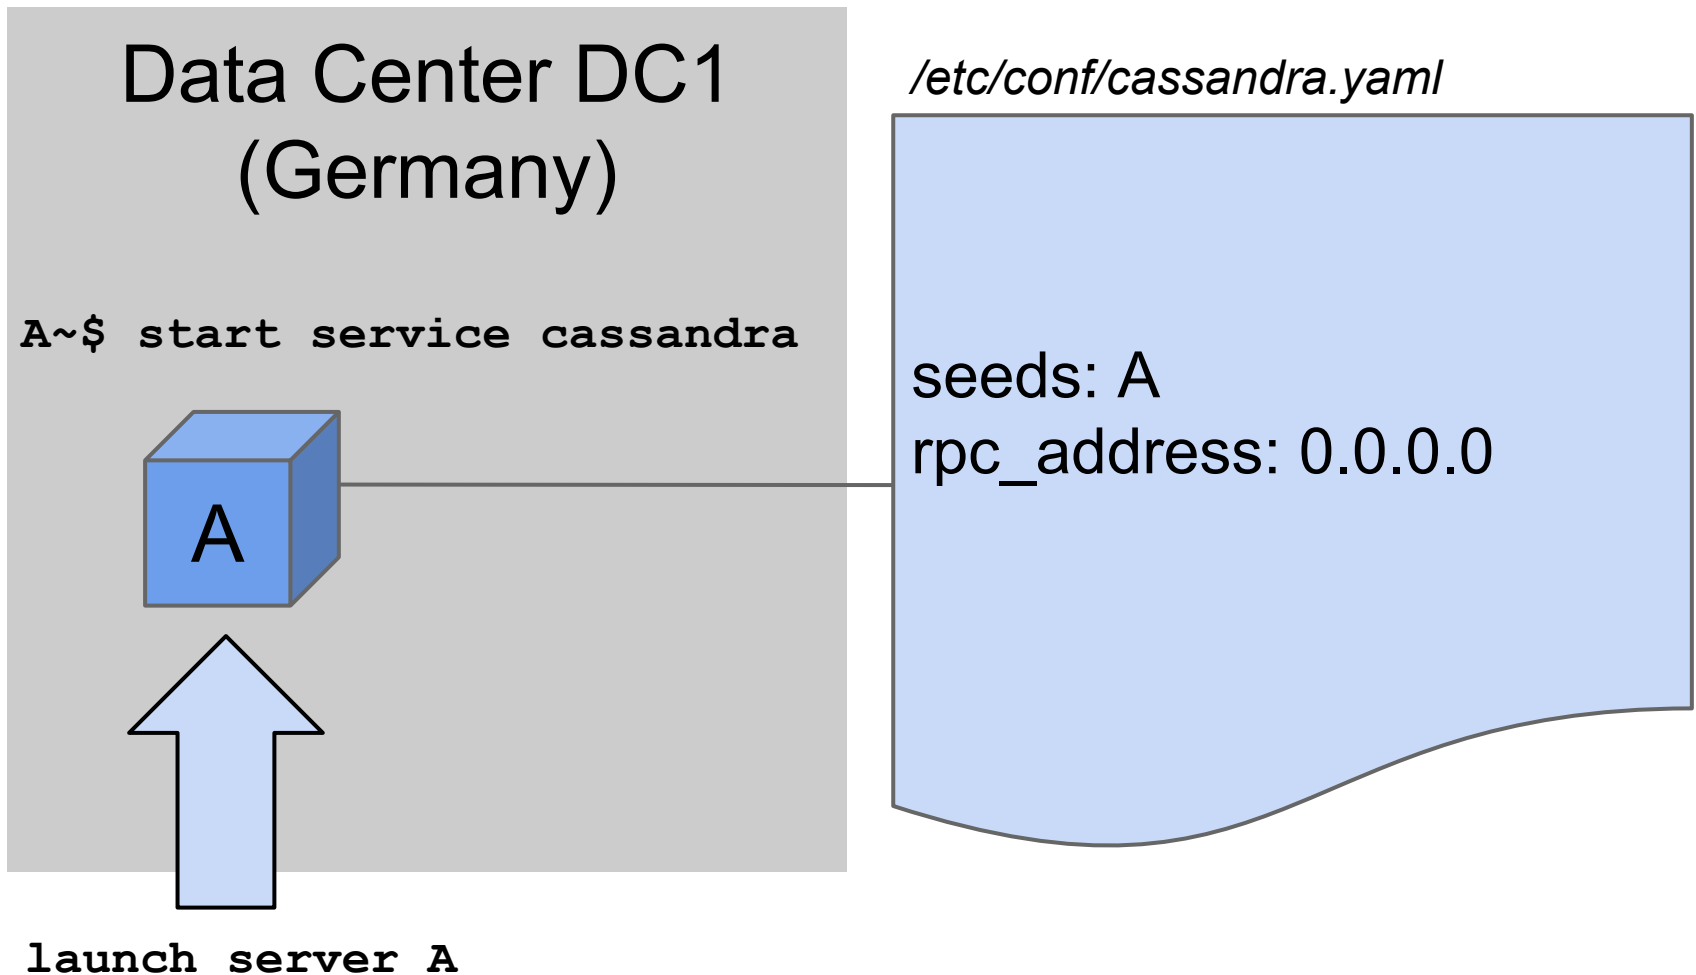
Follows Table

user_id	follows_list
104	[101,117]

Followed Table

id	followed_list
101	[104,109]


Exercise: Launch 1 Cassandra Node



Exercise: Start CQL Shell

Data Center DC1
(Germany)

```
A~$ cqlsh
```



Intro: CLI, CQL2, CQL3

- CQL is “SQL for Cassandra”
- Cassandra CLI deprecated, CQL2 deprecated
- CQL3 is default since Cassandra 1.2

```
$ cqlsh
```

- Pipe scripts into cqlsh

```
$ cat cql_script | cqlsh
```

- Source files inside cqlsh

```
cqlsh> SOURCE '~/cassandra_training/cql3/  
01_create_keyspaces';
```

CQL3

- Create a keyspace
- Create a column family
- Insert data
- Alter schema
- Update data
- Delete data
- Apply batch operation
- Read data
- Secondary Index
- Compound Primary Key
- Collections
- Consistency level
- Time-To-Live (TTL)
- Counter columns
- sstable2json utility tool

Create a SimpleStrategy keyspace

- Create a keyspace with SimpleStrategy and "replication_factor" option with value "3" like this:

```
cqlsh> CREATE KEYSPACE <ksname>
      WITH REPLICATION =
      {'class':'SimpleStrategy',
      'replication_factor':3};
```


Exercise: Create a SimpleStrategy keyspace

- Create a keyspace "simplifiedb" with SimpleStrategy and replication factor 1.

Exercise: Create a SimpleStrategy keyspace

```
cqlsh> CREATE KEYSPACE simpledb
        WITH REPLICATION = {
            'class' : 'SimpleStrategy',
            'replication_factor' : 1 };
cqlsh> DESCRIBE KEYSPACE simpledb;
```

Create a NetworkTopologyStrategy keyspace

Create a keyspace with NetworkTopologyStrategy and strategy option "DC1" with a value of "1" and "DC2" with a value of "2" like this:

```
cqlsh> CREATE KEYSPACE <ksname>
        WITH REPLICATION = {
            'class': 'NetworkTopologyStrategy',
            'DC1': 1,
            'DC2': 2
        };
```

Exercise Create Table “users”

- Connect to the "twotter" keyspace.

```
cqlsh> USE twotter;
```

- Create new column family (Table) named "users".

```
cqlsh:twotter> CREATE TABLE users (  
                                id int PRIMARY KEY,  
                                name text,  
                                email text  
                                );
```

```
cqlsh:twotter> DESCRIBE TABLES;
```

```
cqlsh:twotter> DESCRIBE TABLE users;
```

*we use int instead of uuid in the exercises for the sake of readability

Exercise: Create Table "messages"

- Create a new Table named "messages" with the attributes "posted_on", "user_id", "user_name", "body", and a primary key that consists of "user_id" and "posted_on".

```
cqlsh:twotter> CREATE TABLE messages (  
    posted_on bigint,  
    user_id int,  
    user_name text,  
    body text,  
    PRIMARY KEY (user_id, posted_on)  
);
```

*we use bigint instead of timeuuid in the exercises for the sake of readability

Exercise: Insert data into Table "users" of keyspace "twotter"

```
cqlsh:twotter>
```

```
INSERT INTO users(id, name, email)
VALUES (101, 'otto', 'otto@abc.de');
```

```
cqlsh:twotter> ... insert more records ...
```

```
cqlsh> SOURCE
'~/cassandra_training/cql3/03_insert';
```

Exercise: Insert message records

```
cqlsh:twotter>
```

```
INSERT INTO messages (user_id, posted_on,  
    user_name, body)  
VALUES (101, 1384895178, 'otto', 'Hello!');
```

```
cqlsh:twotter> SELECT * FROM messages;
```

Read data

```
cqlsh:twotter> SELECT * FROM users;
```

id	email	name
105	g@rd.de	gerd
104	linda@abc.de	linda
102	null	jane
106	heinz@xyz.de	heinz
101	otto@abc.de	otto
103	null	karl

Update data

```
cqlsh:twotter> UPDATE users
                  SET email = 'jane@smith.org'
                  WHERE id = 102;
```

id	email	name
105	g@rd.de	gerd
104	linda@abc.de	linda
102	jane@smith.org	jane
106	heinz@xyz.de	heinz
101	otto@abc.de	otto
103	null	karl

Delete data

- Delete columns

```
cqlsh:twotter> DELETE email  
                  FROM users  
                  WHERE id = 105;
```

- Delete an entire row

```
cqlsh:twotter> DELETE FROM users  
                  WHERE id = 106;
```

Delete data

id	email	name
105	null	gerd
104	linda@abc.de	linda
102	jane@smith.org	jane
101	otto@abc.de	otto
103	null	karl

Batch operation

- Execute multiple mutations with a single operation

```
cqlsh:twotter>
```

```
BEGIN BATCH
```

```
    INSERT INTO users(id, name, email)
```

```
        VALUES(107, 'john', 'j@doe.net')
```

```
    INSERT INTO users(id, name)
```

```
        VALUES(108, 'michael')
```

```
    UPDATE users
```

```
        SET email = 'michael@abc.de'
```

```
        WHERE id = 108
```

```
        DELETE FROM users WHERE id = 105
```

```
APPLY BATCH;
```

Batch operation

id	email	name
107	j@doe.net	john
108	michael@abc.de	michael
104	linda@abc.de	linda
102	jane@smith.org	jane
101	otto@abc.de	otto
103	null	karl

Secondary Index

```
cqlsh:twotter>
```

```
CREATE INDEX name_index ON users(name);
```

```
cqlsh:twotter>
```

```
CREATE INDEX email_index ON users(email);
```

```
cqlsh:twotter> SELECT name, email FROM  
    users WHERE name = 'otto';
```

```
cqlsh:twotter> SELECT name, email FROM  
    users WHERE email = 'michael@abc.de';
```

Alter Table Schema

```
cqlsh:twotter>
```

```
ALTER TABLE users ADD password text;
```

```
cqlsh:twotter>
```

```
ALTER TABLE users  
    ADD password_reset_token text;
```

* Given its flexible schema, Cassandra's CQL ALTER finishes much quicker than RDBMS SQL ALTER where all existing records need to be updated.

Alter Table Schema

id	email	name	password	password_reset_token
107	j@doe.net	john	null	null
108	michael@abc.de	michael	null	null
104	linda@abc.de	linda	null	null
102	jane@smith.org	jane	null	null
101	otto@abc.de	otto	null	null
103	null	karl	null	null

Collections - Set

CQL3 introduces **collections** for storing complex data structures, namely the following: **set**, **list**, and **map**. This is the CQL way of modelling many-to-one relationships.

1. Let us add a set of "hobbies" to the Table "users".

```
cqlsh:twotter> ALTER TABLE users ADD hobbies  
    set<text>;  
cqlsh:twotter> UPDATE users SET hobbies =  
    hobbies +  
    {'badminton', 'jazz'} WHERE id = 101;
```

Collections - List

2. Now create a Table "followers" with a list of followers.

```
cqlsh:twotter> CREATE TABLE followers (  
    user_id int PRIMARY KEY,  
    followers list<text>);
```

```
cqlsh:twotter> INSERT INTO followers (  
    user_id, followers)  
VALUES (101, ['willi', 'heinz']);
```

```
cqlsh:twotter> SELECT * FROM followers;  
    user_id | followers
```

```
-----+-----
```

```
101      | ['willi', 'heinz']
```

Collections - Map

3. Add a map to the Table "messages".

```
cqlsh:twotter>
```

```
ALTER TABLE messages
```

```
    ADD comments map<text, text>;
```

```
cqlsh:twotter>
```

```
UPDATE messages
```

```
    SET comments = comments + {'otto':'thx!'}
```

```
    WHERE user_id = 103
```

```
    AND posted_on = 1384895223;
```

Consistency Level

- Set the consistency level for all subsequent requests:

```
cqlsh:twotter> CONSISTENCY ONE;  
cqlsh:twotter> CONSISTENCY QUORUM;  
cqlsh:twotter> CONSISTENCY ALL;
```

- Show the current consistency level:

```
cqlsh:twotter> CONSISTENCY;
```

Exercise: Consistency Level

- Set the consistency level to ANY and execute a SELECT statement.

Exercise: Consistency Level

- Set the consistency level to ANY and execute a SELECT statement.

Bad Request: ANY ConsistencyLevel is only supported for writes

Exercise: Time-To-Live (TTL)

- Insert a user record with a password reset token with a 77 second TTL value.

```
cqlsh:twotter>
```

```
INSERT INTO users (id, name,  
    password_reset_token)  
VALUES (109, 'timo', 'abc-xyz-123')  
USING TTL 77;
```

Exercise: Time-To-Live (TTL)

- The INSERT statement before will delete the entire user record after 77 seconds.
- This is what we actually wanted to do:

```
cqlsh:twotter> INSERT INTO users (id, name)
                VALUES (110, 'anna');
cqlsh:twotter> UPDATE users USING TTL 77
                SET password_reset_token =
                    'abc-xyz-123'
                WHERE id = 110;
```


Time-To-Live (TTL)

- Check the TTL expiration time in seconds.

```
cqlsh:twotter> SELECT TTL  
                  (password_reset_token)  
                  FROM messages  
                  WHERE user_id = 110;
```

Counter Columns

Create a Counter Column Table that counts "upvote" and "downvote" events.

```
cqlsh:twotter> CREATE TABLE votes (  
    user_id int,  
    msg_created_on bigint,  
    upvote counter,  
    downvote counter,  
    PRIMARY KEY (user_id, msg_created_on)  
);
```

Counter Columns

```
cqlsh:twotter>
```

```
UPDATE votes SET upvote = upvote + 1  
  WHERE user_id = 101  
  AND msg_created_on = 1234;
```


```
cqlsh:twotter>
```

```
UPDATE votes  
  SET downvote = downvote + 2  
  WHERE user_id = 101  
  AND msg_created_on = 1234;
```

```
cqlsh:twotter> SELECT * FROM votes;
```

sstable2json utility tool

```
$ sstable2json  
var/lib/cassandra/data/twotter/users/*.db >  
*.json
```



Exercise: sstable2json

- Insert a few records
- Flush the *users* column family to disk and create a json representation of a *.db file.

Solution: sstable2json

```
$ nodetool flush twotter users
$ sstable2json
/var/lib/cassandra/data/twotter/users/twotter-users-jb-1-Data.db > twotter-users.json
$ cat twotter-users.json
[{"key": "000000069", "columns": [[ "", "",
1384963716697000], ["email", "g@rd.de",
1384963716697000], ["name", "gerd",
1384963716697000]] },
{"key": "000000068", "columns": [[ "", "",
1384963716685000], ...
```

CQL v3.1.0

(New in Cassandra 2.0)

- IF keyword
- Lightweight transactions (“Compare-And-Set”)
- Triggers (experimental!!)
- CQL paging support
- Drop column support
- SELECT column aliases
- Conditional DDL
- Index enhancements
- cqlsh COPY

IF Keyword

```
cqlsh> DROP KEYSPACE twotter;
```

```
cqlsh> DROP KEYSPACE twotter;
```

**Bad Request: Cannot drop non existing
keyspace 'twotter'.**

```
cqlsh> DROP KEYSPACE IF EXISTS twotter;
```


Lightweight Transactions

- Compare And Set (CAS)
- Example: without CAS, two users attempting to create a unique user account in the same cluster could overwrite each other's work with neither user knowing about it.

Lightweight Transactions

1. Register a new user

```
cqlsh:twotter>
```

```
INSERT INTO users (id, name, email)  
VALUES (110, 'franz', 'fr@nz.de')  
IF NOT EXISTS;
```

2. Perform a CAS reset of Karl's email.

```
cqlsh:twotter>
```

```
UPDATE users  
SET email = 'franz@gmail.com'  
WHERE id = 110  
IF email = 'fr@nz.de';
```

Exercise: Lightweight Transactions

- Perform a failing CAS e-mail reset:

```
cqlsh:twotter>
```

```
UPDATE users
```

```
    SET email = 'franz@ABC.de'
```

```
...
```

Exercise: Lightweight Transactions

- Perform a failing CAS e-mail reset:

```
cqlsh:twotter>
```

```
UPDATE users  
  SET email = 'franz@ABC.de'  
  WHERE id = 110  
  IF email = 'fr@nz.de';
```

```
[applied] | email
```

```
-----+-----
```

```
False    | franz@gmail.com
```

Exercise: Lightweight Transactions

- Write a password reset method by using an expiring password_reset_token column and a CAS password update query.

Exercise: Lightweight Transactions

```
cqlsh:twotter>
```

```
UPDATE users USING TTL 77
```

```
    SET password_reset_token = 'abc-xyz-123'
```

```
    WHERE id = 110;
```

```
cqlsh:twotter>
```

```
UPDATE users
```

```
    SET password = 'geheim!'
```

```
    WHERE id = 110
```

```
    IF password_reset_token = 'abc-xyz-123';
```

Create a Trigger (experimental feature)

- Triggers are written in Java.
- Triggers are currently an experimental feature in Cassandra 2.0. Use with caution!

```
cqlsh:twotter>
```

```
CREATE TRIGGER myTrigger ON users USING 'org.  
apache.cassandra.triggers.InvertedIndex'
```